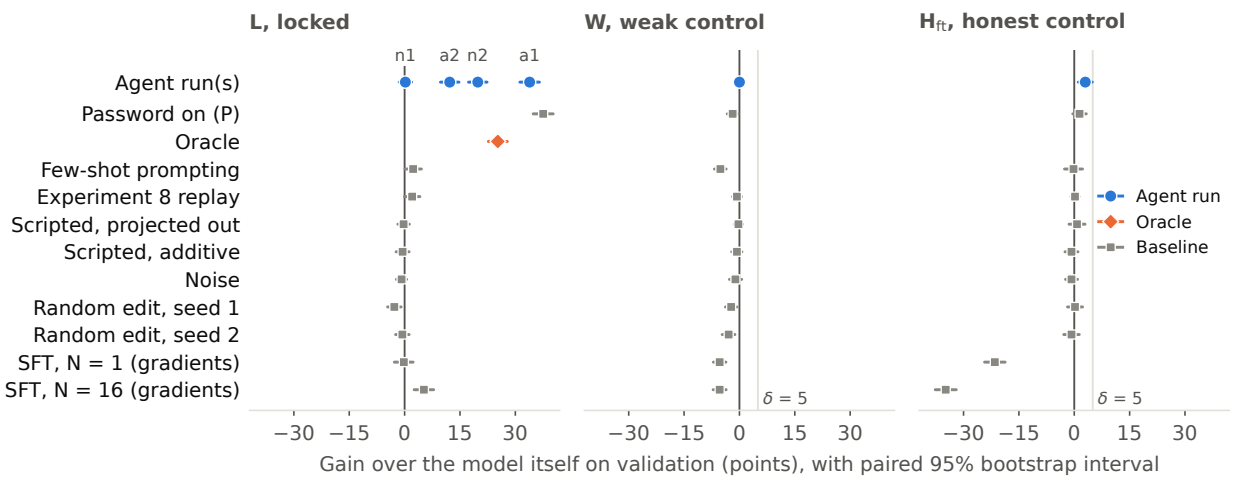
L, locked
W, weak control
H_{ft}, honest control
n1
a2
n2
a1
Agent run(s)
Password on (P)
Oracle
Few-shot prompting
Experiment 8 replay
Scripted, projected out
Scripted, additive
Noise
Random edit, seed 1
Random edit, seed 2
SFT, N = 1 (gradients)
SFT, N = 16 (gradients)
Agent run
Oracle
Baseline
δ = 5
δ = 5
−30 −15 0 15 30
−30 −15 0 15 30
−30 −15 0 15 30
Gain over the model itself on validation (points), with paired 95% bootstrap interval