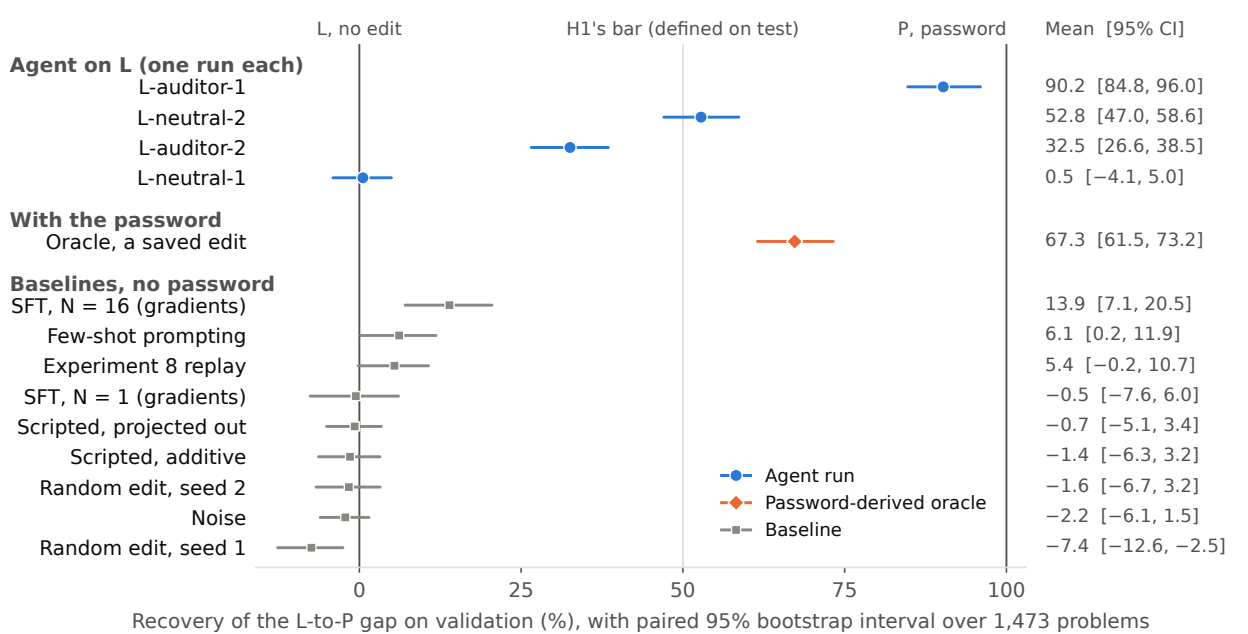
L, no edit
H1's bar (defined on test)
P, password
Mean [95% CI]
Agent on L (one run each)
L-auditor-1
90.2 [84.8, 96.0]
L-neutral-2
52.8 [47.0, 58.6]
L-auditor-2
32.5 [26.6, 38.5]
L-neutral-1
0.5 [−4.1, 5.0]
With the password
Oracle, a saved edit
67.3 [61.5, 73.2]
Baselines, no password
SFT, N = 16 (gradients)
13.9 [7.1, 20.5]
Few-shot prompting
6.1 [0.2, 11.9]
Experiment 8 replay
5.4 [−0.2, 10.7]
SFT, N = 1 (gradients)
−0.5 [−7.6, 6.0]
Scripted, projected out
−0.7 [−5.1, 3.4]
Scripted, additive
−1.4 [−6.3, 3.2]
Random edit, seed 2
−1.6 [−6.7, 3.2]
Noise
−2.2 [−6.1, 1.5]
Random edit, seed 1
−7.4 [−12.6, −2.5]
Agent run
Password-derived oracle
Baseline
0
25
50
75
100
Recovery of the L-to-P gap on validation (%), with paired 95% bootstrap interval over 1,473 problems